

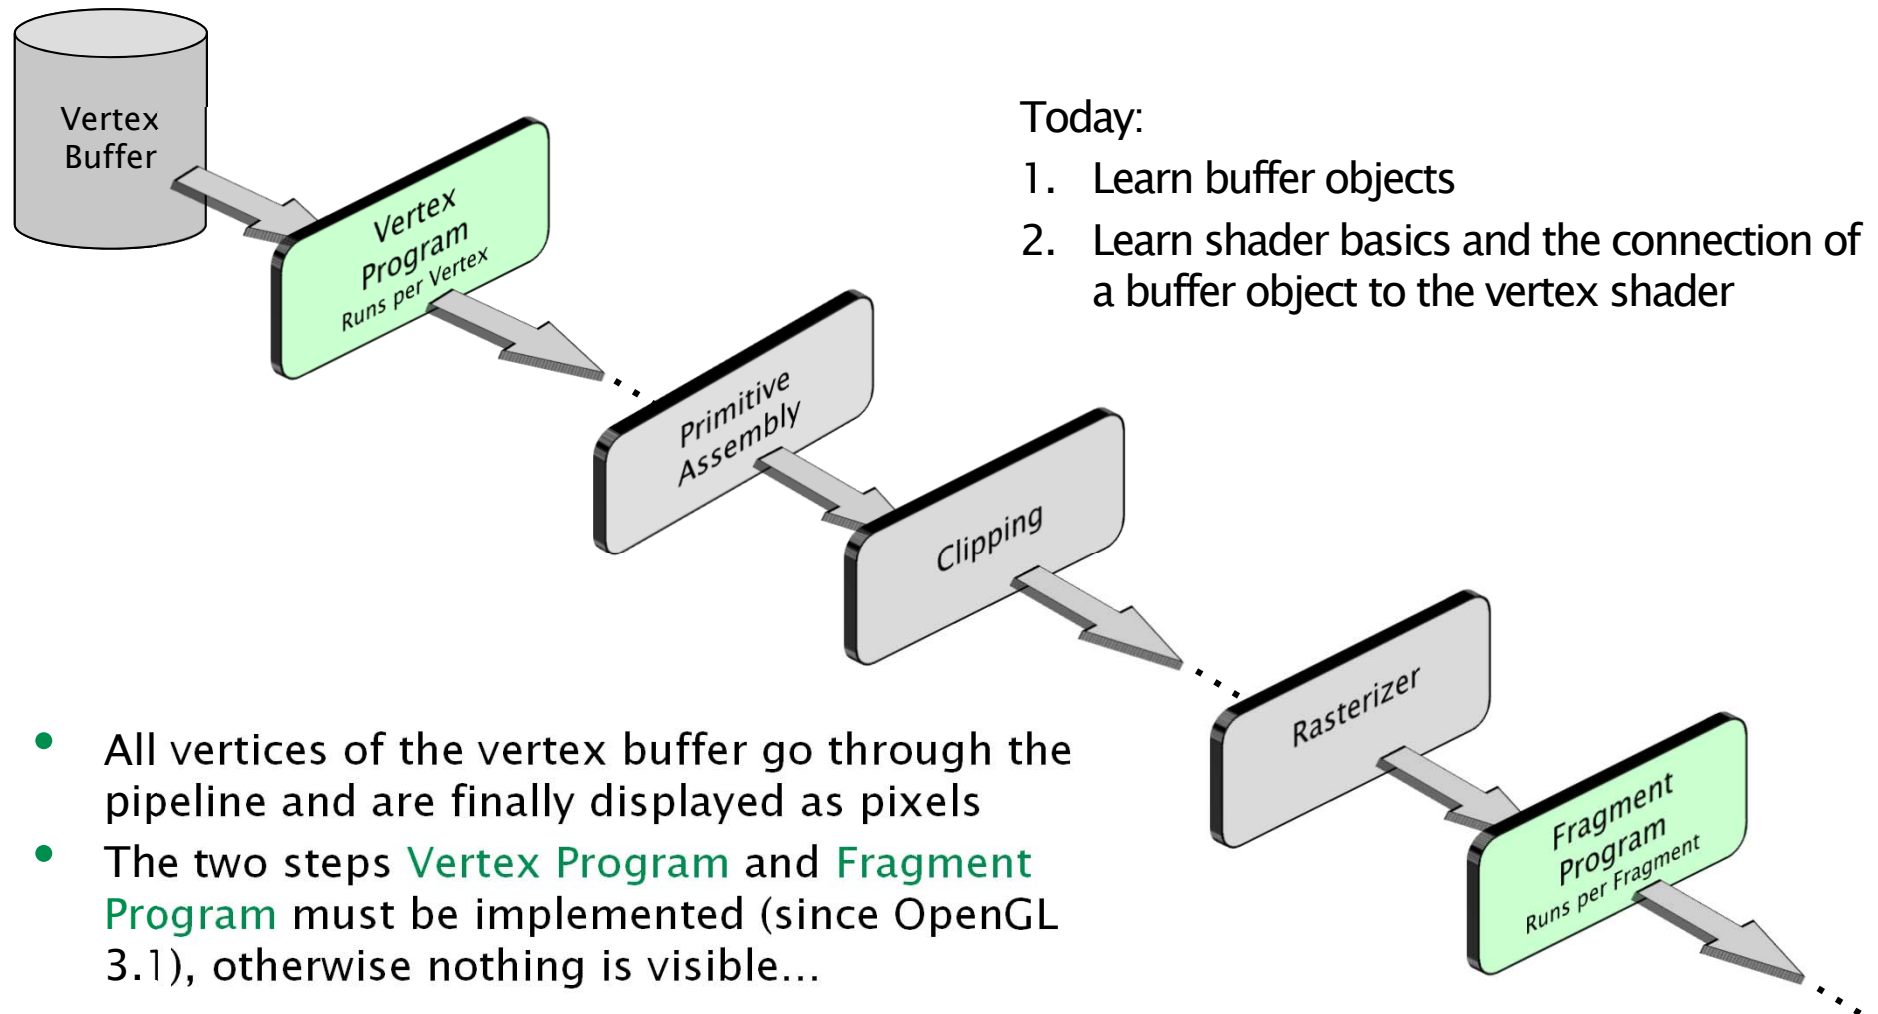
(3) Buffer Objects

GPU Programming
T. Grosch

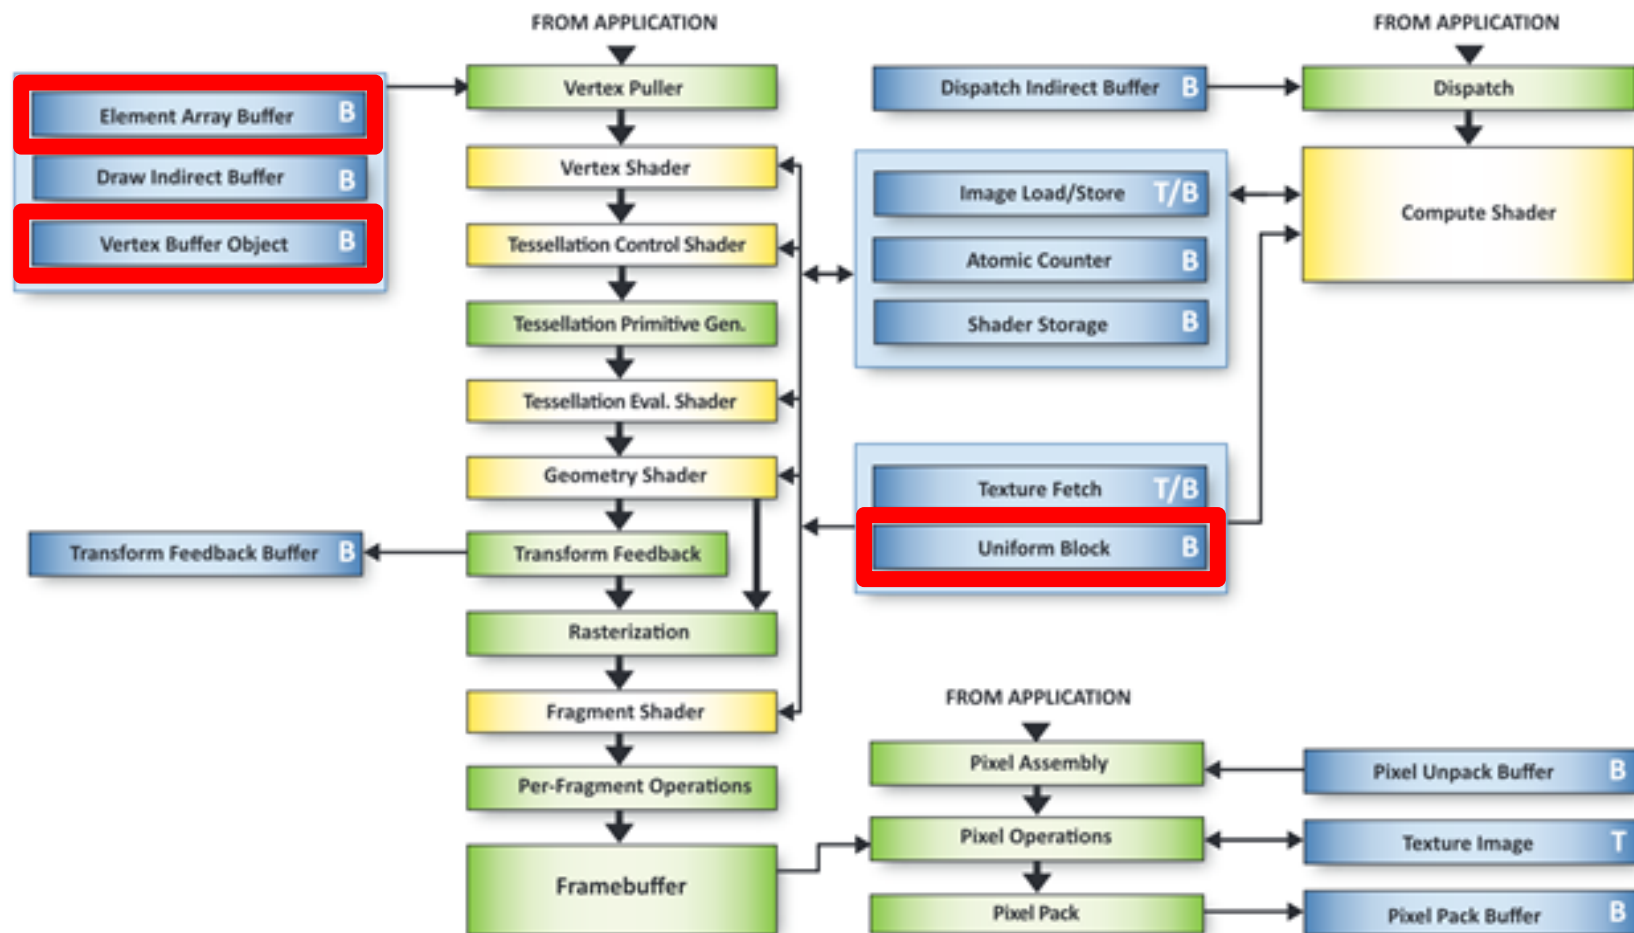
Today

- Buffer Objects
- Repeat Vertex Buffer Objects and Vertex Array Objects (but with a new semantic)
- Uniform Buffer Objects
- Vertex and Fragment Shader Intro

Simplified OpenGL Pipeline



OpenGL 4.4




OpenGL Buffers

- Buffer = Block of data in (GPU) memory with additional description
 - OpenGL uses buffers at several places, not only as vertex data
 - Examples:
 - GL_ARRAY_BUFFER : Vertex Data
 - GL_ELEMENT_ARRAY_BUFFER : Index Data
 - PIXEL_(UN)PACK_BUFFER : RGB Color Data
 - GL_TRANSFORM_FEEDBACK_BUFFER : Vertex Data
 - GL_SHADER_STORAGE_BUFFER : Shader Input/Output Data
 - GL_UNIFORM_BUFFER : Shader Input Data
 - ...


Vertex Buffer Object

- We start with the Vertex Buffer Object
- Data = xyz positions of the vertices
- Goal: Draw a mesh consisting of graphical primitives
 - Triangles, Lines, Points, ...
- Optional: Additional vertex information
 - Vertex Normals
 - Vertex Colors
 - Vertex Texture Coordinates
 - ...


Primitive Types in OpenGL (Excerpt)




GL_POINTS



GL_LINE_STRIP




GL_LINE_LOOP




GL_POLYGON


Multiple Primitives



GL_LINES




GL_TRIANGLES



GL_QUADS

Rasterization using OpenGL

- The default coordinates for vertices are not pixel coordinates !
 - The standard coordinate system in OpenGL ranges from -1 to $+1$ in both directions
- Definition:
 - The x axis points to the right
 - The y axis points upward
 - Attention: In GLFW, the y axis points down !
 - The origin is in the center of the image
- All objects with coordinates outside the range $[-1, +1]$ are not visible !



Drawing with OpenGL

1. The complete vertex data is transferred (as a block) from CPU to GPU and remains there (usually).
2. A description for the vertex data is added, which describes the memory layout (start address, number of vertices, how many coordinates 2D/3D, float or int as base type, ...)
3. Using a simple draw-command, the whole block can be drawn (without any further data transmission)

This type of geometry data is called a **Vertex Buffer Object (VBO)**

Vertex Buffer Checklist

- Define the vertex coordinates, colors, normals... (as C arrays)
- Generate Buffer Object(s) `glGenBuffer`, `glBindBuffer`
- Allocate memory `glBufferStorage`
- Copy the vertex coordinates data `glBufferSubData`
- Generate a Vertex Array Object `glGenVertexArrays` and `glBindVertexArray`
- Describe and enable the vertex attributes
`glVertexAttribFormat`, `glEnableVertexAttribArray`
- Set the bindings for the attributes `glVertexAttribBinding`,
`glBindVertexBuffer`
- Define the attribute numbers in the vertex shader
(`layout(location = ...)`)
- Draw the geometry `glDrawArrays`

Description of the Commands

- We will go through all these commands and their parameters (which is quite a lot...)
- Some of them are related to the input of the vertex shader (we assume that there is a vertex shader and explain this later)
- Afterwards we will have a look at a small example to understand the connection between the commands
 - The commands are slightly different to those used in Computer Graphics 1 for an older OpenGL version

Vertex Buffer Object Init

- OpenGL reserves IDs for different VBOs, such an ID can be generated with:

```
GLuint vboID;  
glGenBuffers( 1 , &vboID );
```

- Afterwards we can use

```
glBindBuffer( GL_ARRAY_BUFFER, vboID );
```

to bind the buffer for upcoming buffer related commands

Allocating Memory for the VBO

- Assigning vertex data to a VBO is implemented with

```
glBufferStorage(target, size, data, flags);
```

- `target` is set to `GL_ARRAY_BUFFER` (for now)
 - `size` is the size in bytes
 - `data` is the address in main memory
 - `flags` is set to `GL_DYNAMIC_STORAGE_BIT` or none (0) (for now).
- Using this command, the vertex data is transferred from CPU to GPU. In case of `data == NULL`, only the GPU memory is allocated.

Assigning Partial Vertex Data to a VBO

- (Re-)assigning (partial) vertex data to an existing VBO is implemented with

```
glBufferSubData(target, offset, size, data);
```

- `target` is set to `GL_ARRAY_BUFFER` (for now)
- `offset` is added to the buffer start address and decides where the `data` should be copied to (measured in bytes). Can be set to 0 if the whole image should be copied.
- `size` is the size in bytes
- `data` is the address in main memory

Vertex Array Object

- In addition to the Vertex Buffer Object(s), we need a **Vertex Array Object (VAO)** to define how the buffer data is streamed into the vertex shader
- VAOs come in two different flavors:
 - Aggregated buffer bindings (old)
One VAO per object required
 - Separated: only specifies the format
One VAO per different vertex shader input
- To create and use a VAO call `glGenVertexArrays` and `glBindVertexArray`

Vertex Array Object Init

- OpenGL reserves IDs for different VAOs, such an ID can be generated with:

```
GLuint vaoID;  
glGenVertexArrays( 1 , &vaoID );
```

- Afterward we can use

```
glBindVertexArray( vaoID );
```

to define (bind) the current VAO. The following commands are then related to this VAO.

Detailed Data Description

- Up to now, we only know the total size of the geometry data, but we do not know the **memory layout** of the vertex data. Therefore, we need an additional format description:

```
glVertexAttribFormat(attributeIndex, size, type,  
                      normalized, offset);
```

- **attributeIndex**: Attribute number in the vertex shader (later)
- **size**: number of components per vertex (2 for xy, 3 for xyz, ...)
- **type**: Base type, typically `GL_FLOAT` or `GL_UNSIGNED_BYTE`
- **normalized**: Integer types are converted to $[-1,1]$ or $[0,1]$ values (we set this to `GL_FALSE` for now)
- **offset**: byte offset within each vertex (data can be interleaved: e.g. positions and normals in a single buffer). Can be set to 0 for simple (non-interleaved) buffers.

Activate Vertex Attribute

- Each vertex attribute must be activated:

```
glEnableVertexAttribArray (attribIndex) ;
```

otherwise it does not work (!)

- Typical vertex attributes are
 - Position, normal, color, texture coordinate, ...

Bind a Vertex Buffer

- To use a buffer as vertex buffer for the draw call we need to bind it to a so-called binding index:

```
glBindVertexBuffer( bindingIndex, vboID,  
                    offset , stride );
```

- **bindingIndex**: Binding index of the vertex buffer (this is not the attribute index!)
- **vboID**: buffer object ID to be assigned to the binding index
- **stride**: Gap between two successive elements in memory (in bytes)
- **offset**: is added to the address given by `glBufferStorage`. Can be set to zero in most cases.

Note: In Computer Graphics 1, we use `stride = 0` for tightly packed vertices, because the effective stride is computed automatically. This does no longer work here !

Vertex Attribute Binding

- Connect a VBO (given by the binding index) to the input of the vertex shader (given by the attribute index)
- Which attribute is taken from which buffer:

```
glVertexAttribBinding( attribIndex,  
                       bindingIndex );
```

- **attribIndex**: The attribute location in the shader
- **bindingIndex**: The vertex buffer from which the attribute is taken
- Must be called after **glVertexAttribFormat**

Draw the VBO (finally...)

- After defining all the VBO data, we can draw the VBO using

```
glDrawArrays(mode, first, count);
```

- Here `mode` is the geometry type, e.g. `GL_TRIANGLES` or `GL_LINES`.
- Draw `count` vertices starting from index `first`. The vertex coordinates are not necessary here.

Clean up

- VAO and VBO can be deleted with

```
glDeleteBuffers (...);
```

```
glDeleteVertexArrays (...);
```

Example

- To explain all these commands, we use geometry with **positions, normals and texture coordinates** (3 attributes)
- We assume that we have two triangles = 6 vertices
- For positions and normals, we use 3 floats for xyz coordinates (=12 bytes each)
- For texture coordinates we use 2 floats for st coordinates (= 8 bytes each)
- We store positions and normals **interleaved** in one buffer object
 - Like this: position1 normal1 position2 normal2 ...
- Texture coordinates are stored in a second buffer object
 - This is an arbitrary design decision, we could also use three individual buffers or only a single buffer with three attributes interleaved

Define Vertex Coordinates

- The vertex data is defined completely, e.g. as a C array:

```
GLfloat positions_normals[6][6] = {    // interleaved
    {0.0, 0.0, 0.0,    0.0, 0.0, 1.0}, // vertex 1 normal 1
    {1.0, 0.0, 0.0,    0.0, 0.0, 1.0}, // vertex 2 normal 2
    {1.0, 1.0, 0.0,    0.0, 0.0, 1.0}, // vertex 3 normal 3
    {0.0, 0.0, 0.0,    1.0, 0.0, 0.0}, // vertex 4 normal 4
    {0.0, 1.0, 0.0,    1.0, 0.0, 0.0}, // vertex 5 normal 5
    {0.0, 1.0, 1.0,    1.0, 0.0, 0.0}  // vertex 6 normal 6
};

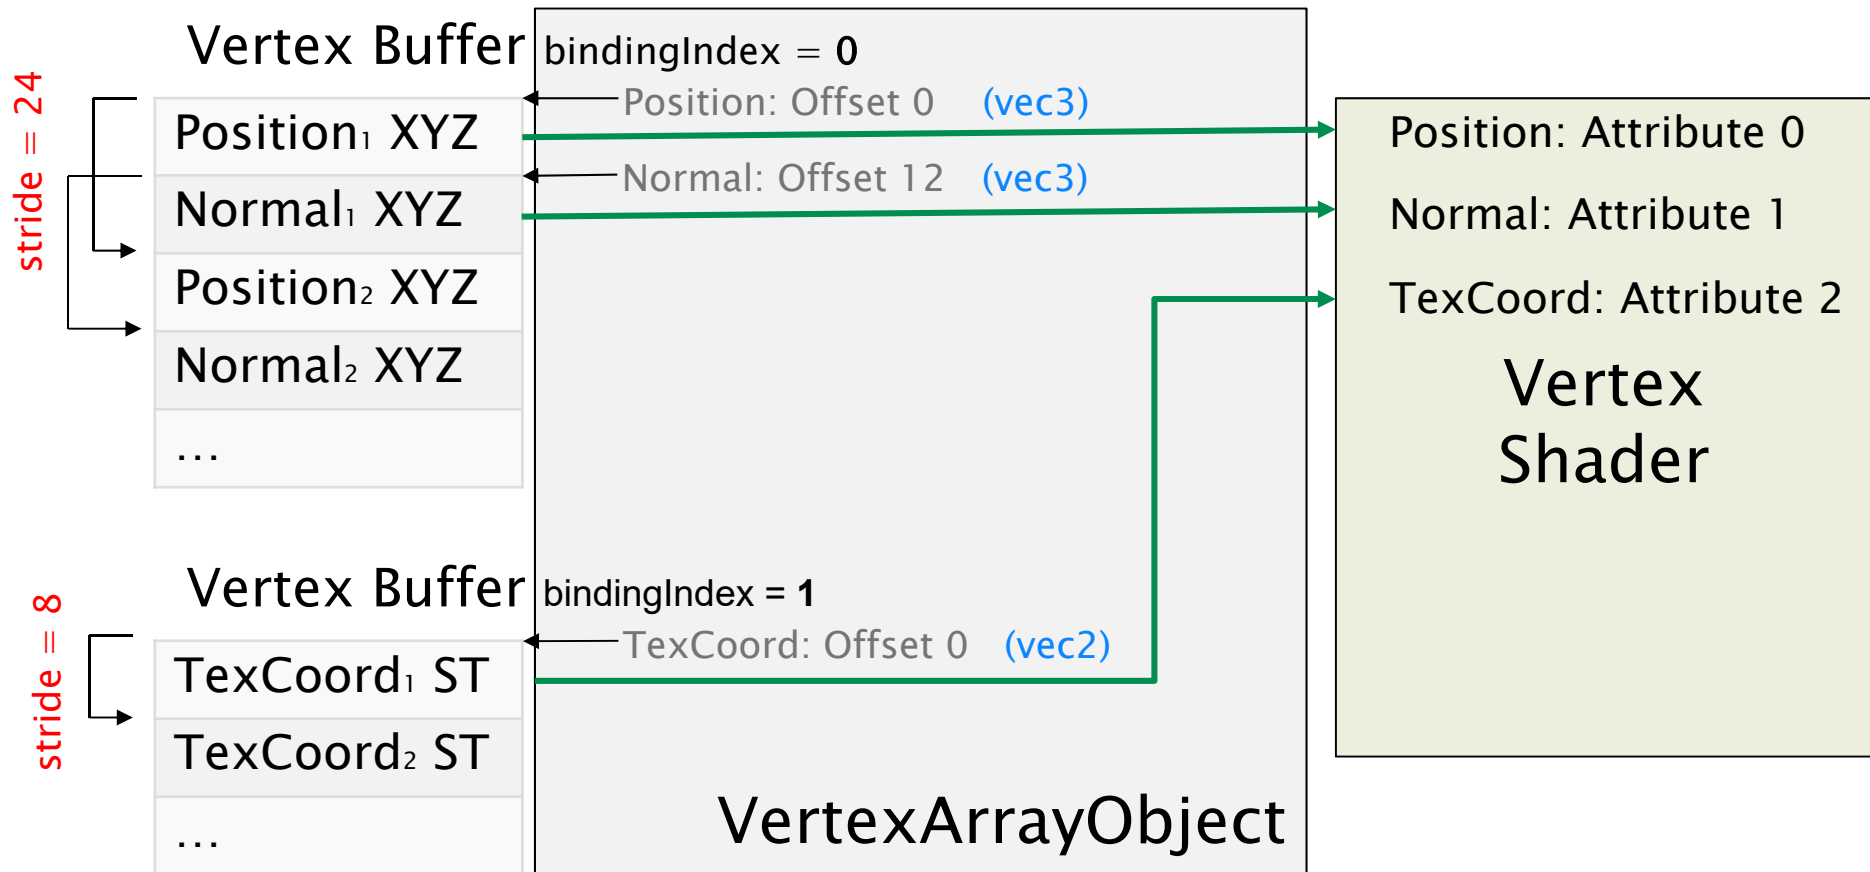
GLfloat texcoords[6][2] = { // 6 texcoords
    {0.0, 0.0}, {0.0, 1.0}, {1.0, 0.0},
    {0.0, 1.0}, {0.0, 0.0}, {1.0, 0.0}
};
```



Example: Create two VBOs

```
GLuint vboPosNormal, vboTexCoord;  
GLfloat positions_normals[6][6] = { ... };  
GLfloat texcoords[6][2] = { ... };  
  
glGenBuffers(1, &vboPosNormal);  
glBindBuffer(GL_ARRAY_BUFFER, vboPosNormal);  
glBufferStorage(GL_ARRAY_BUFFER, 6 * 6 * sizeof(float),  
                positions_normals, 0);  
  
glGenBuffers(1, &vboTexCoord);  
glBindBuffer(GL_ARRAY_BUFFER, vboTexCoord);  
glBufferStorage(GL_ARRAY_BUFFER, 6 * 2 * sizeof(float),  
                texcoords, 0);
```

Example: VAO Overview



Example: Create a VAO and set Connections

```
GLuint vao;  
glGenVertexArrays(1, &vao);  
glBindVertexArray(vao);  
  
glEnableVertexAttribArray(0);           // position  
glVertexAttribFormat(0, 3, GL_FLOAT, GL_FALSE, 0);  
glVertexAttribBinding(0, 0);  
  
glEnableVertexAttribArray(1);          // normal  
glVertexAttribFormat(1, 3, GL_FLOAT, GL_FALSE, 12);  
glVertexAttribBinding(1, 0);  
  
glEnableVertexAttribArray(2);          // texture coords  
glVertexAttribFormat(2, 2, GL_FLOAT, GL_FALSE, 0);  
glVertexAttribBinding(2, 1);
```

Example: Draw the VBO

```
void display()  
{  
    glClear(...);  
    glBindVertexArray(vao);  
    glBindVertexBuffer(0, vboPosNormal, 0, 24);  
    glBindVertexBuffer(1, vboTexCoord, 0, 8);  
    glDrawArrays(GL_TRIANGLES, 0, 6);  
}
```


Note that there are redundancies in OpenGL:

For example, we assign the binding index 0 to the VBO `vboPosNormal`, although `vboPosNormal` is already a unique ID for this VBO...

But for the connection with the shader (`glVertexAttribBinding`) we need the binding index.

Mesh Variants: Indexed VBOs

- Geometry is often described as a connected **Mesh**, where one vertex belongs to multiple triangles
- Up to now, we repeat the vertices → increased memory requirement for multiple points
- Therefore, we can use indexing for VAOs:
 - First we use non-connected list of all vertices
 - Furthermore we use an index-list which describes which vertices should be connected to triangles
- In the example, we have 6 vertices with indices instead of $4 \times 3 = 12$ vertices → less memory
 - One vertex needs $2 \times 4 = 8$ bytes
 - 12 vertices need $12 \times 8 = 96$ bytes
 - three indices need $3 \times 2 = 6$ bytes (short index 16 bit)
 - 6 vertices with 4 triangle index list need only
 $6 \times 8 + 4 \times 6 = 72$ bytes



| Vertex List | Index List |
|---------------------------------|------------|
| x ₀ , y ₀ | 0, 1, 3 |
| x ₁ , y ₁ | 0, 3, 2 |
| x ₂ , y ₂ | 2, 3, 5 |
| x ₃ , y ₃ | 2, 5, 4 |
| x ₄ , y ₄ | |
| x ₅ , y ₅ | |

Assign Index Data

- In addition to the VBO for the vertex data, another buffer for the index data is created using

```
glGenBuffers(1, &ibo);  
glBindBuffer(GL_ELEMENT_ARRAY_BUFFER, ibo);  
glBufferStorage(GL_ELEMENT_ARRAY_BUFFER, size, data, 0);
```

- the index data is marked with `GL_ELEMENT_ARRAY_BUFFER` (instead of `GL_ARRAY_BUFFER`)
 - (no good names here in OpenGL: `GL_VERTEX_BUFFER` and `GL_INDEX_BUFFER` would be better...)
- Again, `data` is the address in main memory and `size` is the size in bytes.

Draw VBO with Indices

- After defining vertex and index data, we can draw with

```
glBindBuffer(GL_ELEMENT_ARRAY_BUFFER, vboIndex);  
glDrawElements(mode, indicesSize, indicesType, offset);
```

- `mode` defines the type of geometry, e.g. `GL_TRIANGLES` or `GL_LINES`.
- `indicesSize` is the number of indices
- `indicesType` defines the base type of the index data
 - Note that only unsigned integer types can be used here e.g. `GL_UNSIGNED_INT`.
 - **Attention:** The command does not work if a signed type is used, e.g. `GL_INT`
- `offset`: is added to the index address given by `glBufferStorage`. Can be set to zero in most cases.

Remark: The indices (elements) are **not** a vertex attribute and therefore not used in a VAO

Example: Create VBO with Indices

```
void vboIndexedInit()  
{  
    GLuint vbo, ibo;  
    GLfloat vertices[6][2] = { ... };  
    GLuint indices[3*4] = { 0,1,3, 0,3,2, ... };  
  
    // generate VBO for vertices...  
  
    glGenBuffers(1, &ibo);  
    glBindBuffer(GL_ELEMENT_ARRAY_BUFFER, ibo);  
    glBufferStorage(GL_ELEMENT_ARRAY_BUFFER, sizeof(indices),  
        indices, 0);  
  
    // generate VAO, define formats and set all connections...  
}
```



Example: Draw VBO with Index List

```
void display()
{
    glClear(...);
    glBindVertexArray(...);
    glBindVertexBuffer(...);
    glBindBuffer(GL_ELEMENT_ARRAY_BUFFER, ibo);
    glDrawElements(GL_TRIANGLES, 3*4, GL_UNSIGNED_INT, NULL);
}
```

Drawing with modern OpenGL

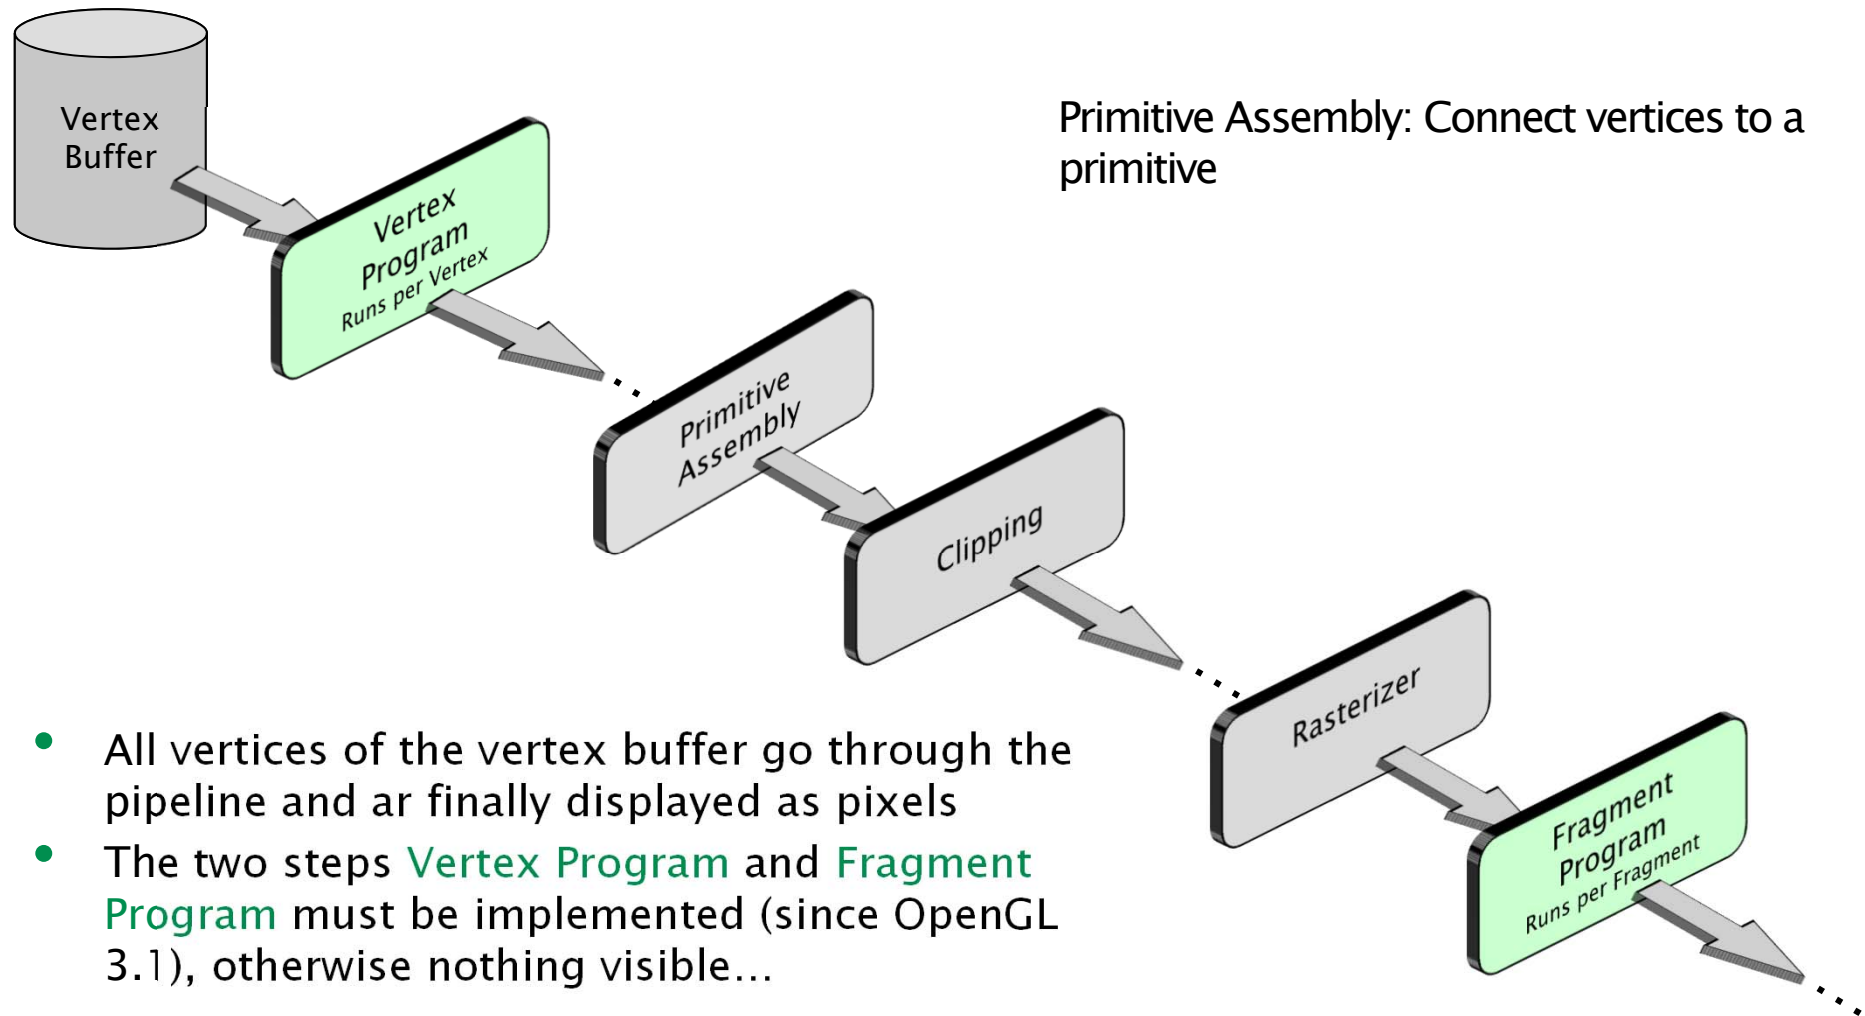
- What happens if we draw the buffer objects using the shown commands → black screen ?
- Although buffer objects are quite complicated we still need more to be able to display an image
- The reason: each vertex goes through a **pipeline** before it is displayed at a pixel coordinate
- Parts of the pipeline must be implemented

Introduction to Vertex and Fragment Shaders



TU Clausthal

Simplified OpenGL Pipeline



Vertex Shader (or Vertex Program)

- ...is called per vertex
 - One vertex is passed to the next pipeline stage
 - The vertex can be modified in the Vertex Shader
 - Transformation, Projectionen, set color → later
 - All other vertices of the geometry are unknown
 - It is not possible to generate new vertices, there is also no (direct) way to remove to vertex
- Simple Vertex Shader
 - pass through

```
#version 440 core
layout (location = 0) in vec4 vPosition; // vertex attribute 0
void main()
{
    gl_Position = vPosition;
}
```

Short Explanation: Vertex Shader

- First Line
 - OpenGL Version number (here: 4.4)
 - **core**: New OpenGL without backwards compatibility
 - (different models have different input/output syntax...)
- **gl_Position**: Output, this vertex is passed to the next stage
- **vPosition**: Input (Vertex position from VBO)
- This is the connection with the VBOs:

```
layout (location = 0) in vec4 vPosition
```

This must fit to the description of the vertex attributes:

```
glVertexAttribFormat(0, 4, GL_FLOAT, GL_FALSE, 0);  
glVertexAttribBinding(0, 0);
```

Fixed Pipeline Stages

- Primitive Assembly
 - After the Vertex Shader, the vertices are assembled to primitives (triangles, lines, ...)
- Clipping
 - The primitives are clipped at the visible range $[-1,1]^2$
 - e.g. Cohen–Sutherland for lines
 - Sutherland–Hodgman for polygons
- Rasterization
 - The visible range of a primitive is subdivided in pixels (fragments). The difference between pixel and fragment will be explained later.

Fragment Shader (or Pixel Shader, Fragment Program)

- The Fragment Shader is called during the rasterization for each generated pixel (fragment)
 - The main task of the Fragment Shader is the definition of the pixel color
 - The Fragment Shader only knows the current pixel, directly reading/writing from/to other pixels is impossible
- Simple example: draw all primitives in red

```
#version 440 core
out vec3 fColor;
void main()
{
    fColor = vec3(1.0, 0.0, 0.0);
}
```


Fragment Shader (or Pixel Shader, Fragment Program)

- Again, the first line sets the version number
- Typically, the Fragment Shader has a single output (the pixel color)
 - the name `vecColor` is arbitrary, using `out` defines it as output
 - In case of multiple outputs, we need to describe the output using `layout(location = ...)` (we will need this later)

Uniform Buffer Objects



TU Clausthal

Uniform Blocks

- Transmit parameters from OpenGL to shader program: **Uniform Variables** (Constants)
- In case of multiple shader programs the number of uniform parameters grows...
- Many of them are often used in multiple shader programs
 - → optimize the access and common usage of uniform variables
- Grouping of multiple uniform Variables in a Block, for example for transformation and projection matrices in a vertex shader:

```
uniform Matrices {  
    mat4 model;  
    mat4 view;  
    mat4 projection;  
};
```



Insert data in Uniform Block

- Create a Buffer Object for the Uniform block using `glGenBuffers`
- To set a Uniform Variable corresponds to the insertion of data in a Uniform Buffer Object using

```
glBindBuffer (...)  
glBufferStorage (...)  
glBufferSubData (...)
```

UBO Generation

- Create Uniform Buffer Object (UBO), bind as current UBO, reserve memory

```
glGenBuffers(1, &myUBO);  
glBindBuffer(GL_UNIFORM_BUFFER, myUBO);  
glBufferStorage(GL_UNIFORM_BUFFER, sizeInBytes, NULL, 0);
```

- Insert Data in Buffer (= set values in Uniform Variables in Block)

```
glBufferSubData(GL_UNIFORM_BUFFER, offset, sizeInBytes,  
                dataPointer);
```

- It is possible to update the entire buffer or a small range using `offset` and `sizeInBytes` (e.g. replacing all matrices or only one)

UBO Connection to Shader

```
glBindBufferBase(GL_UNIFORM_BUFFER,  
                 bindingPoint,  
                 myUBO) ;
```

- This connects the UBO `myUBO` to the binding point `bindingPoint` in the shader
- A binding point in the shader can be defined using

```
layout (binding = ...) uniform ... { };
```

Note that OpenGL is slightly confusing here in comparison to the VAO connection: For Uniform Buffers, we directly assign the UBO ID `myUBO` to a binding point in the shader. In the VAO case, we have to create a binding index for each VBO (which is not in the shader!) and then connect this to an attribute index in the shader.

Uniform Block Example

- Let us look at the matrix example again
- We use the binding point 0 for the uniform block in the vertex shader

```
layout (binding = 0) uniform Matrices {  
    mat4 model;  
    mat4 view;  
    mat4 projection;  
};
```

Uniform Block Example

- Generate a uniform buffer object and reserve memory for three matrices

```
glGenBuffers(1, &myUBO);  
glBindBuffer(GL_UNIFORM_BUFFER, myUBO)  
glBufferStorage(GL_UNIFORM_BUFFER, 3 * sizeof(glm::mat4),  
                NULL, GL_DYNAMIC_STORAGE_BIT);
```

- Fill in the three matrices

```
glm::mat4 modelMatrix, viewMatrix, projMatrix;  
...  
glBufferSubData(GL_UNIFORM_BUFFER, 0 * sizeof(glm::mat4),  
                sizeof(glm::mat4), glm::value_ptr(modelMatrix));  
glBufferSubData(GL_UNIFORM_BUFFER, 1 * sizeof(glm::mat4),  
                sizeof(glm::mat4), glm::value_ptr(viewMatrix));  
glBufferSubData(GL_UNIFORM_BUFFER, 2 * sizeof(glm::mat4),  
                sizeof(glm::mat4), glm::value_ptr(projMatrix));
```


Uniform Block Example

- Connect the UBO with the shader at binding point 0

```
glBindBufferBase(GL_UNIFORM_BUFFER, 0, myUBO);
```

- In the shader we have

```
layout (binding = 0) uniform Matrices {  
    mat4 model;  
    mat4 view;  
    mat4 projection;  
};
```

- Now we can compute things like this in the shader:

```
vec4 viewPos = view * worldPos;
```

Uniform Block Layout

- Memory layout

| Layout qualifier | Description |
|---------------------------|---------------------------|
| <code>shared</code> | Used in multiple programs |
| <code>packed</code> | Optimize memory |
| <code>std140</code> | Standard layout |
| <code>row_major</code> | Matrix row by row |
| <code>column_major</code> | Matrix column by column |

- e.g. `layout(shared, row_major) uniform Matrices { ... };`

That´s all for today

- Next Week: Textures